

Outline (part 1)

- Modeling sequences
- The problem of encoding sequences
- Recurrent Neural Networks (RNNs)

Modeling with Machine Learning: RNN (part 1)




Temporal/sequence problems

- How to cast as a supervised learning problem?



Temporal/sequence problems

- How to cast as a supervised learning problem?



- Historical data can be broken down into feature vectors and target values (sliding window)

$$\begin{bmatrix} 0.82 \\ 0.80 \\ 0.73 \\ 0.72 \end{bmatrix} \quad 0.89$$

$\phi(t) \quad y^{(t)}$



Temporal/sequence problems

- Language modeling: what comes next?

This course has been a tremendous ...



Temporal/sequence problems

- Language modeling: what comes next?

This course has been a tremendous ...

$$\begin{array}{cc} \text{tremendous} & \begin{bmatrix} 0 \\ \vdots \\ 1 \\ 0 \end{bmatrix} \\ & ? \\ \text{a} & \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \\ & \phi(t) \quad y^{(t)} \end{array}$$



Temporal/sequence problems

- Language modeling: what comes next?

This course has been a tremendous ...

$$\begin{array}{cc} \text{a} & \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix} \\ & \text{tremendous} \\ \text{been} & \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix} \\ & \phi(t) \quad y^{(t)} \end{array}$$



What are we missing?


- Sequence prediction problems can be recast in a form amenable to feed-forward neural networks
- But we have to engineer how "history" is mapped to a vector (representation). This vector is then fed into, e.g., a neural network
 - how many steps back should we look at?
 - how to retain important items mentioned far back?
- Instead, we would like to learn how to encode the "history" into a vector



Learning to encode/decode

- Language modeling

This course has been a



success (?)

- Sentiment classification

I have seen better lectures



-1

- Machine translation

I have seen better lectures



Olen nähnyt parempia
luentoja

encoding decoding



Encoding everything

words

$$\begin{bmatrix} .1 \\ .3 \\ .4 \end{bmatrix} \begin{bmatrix} .7 \\ .1 \\ .0 \end{bmatrix} \begin{bmatrix} .2 \\ .8 \\ .3 \end{bmatrix} \dots$$

“Efforts and courage are not enough
without purpose and direction” — JFK

sentences

$$\begin{bmatrix} .2 \\ .3 \\ .6 \end{bmatrix}$$

images

$$\begin{bmatrix} .3 \\ .3 \\ .5 \end{bmatrix}$$



1. Wikimedia, public domain

events

$$\begin{bmatrix} .2 \\ .4 \\ .6 \end{bmatrix}$$




2. Defense.gov, Public Domain,



Example: encoding sentences

- Easy to introduce adjustable “lego pieces” and optimize them for end-to-end performance




<null>

Efforts and courage are not ...



Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance





<null>

Efforts and courage are not ...



Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance





Efforts and courage are not ...



Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance





Efforts and courage are not ...



Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance



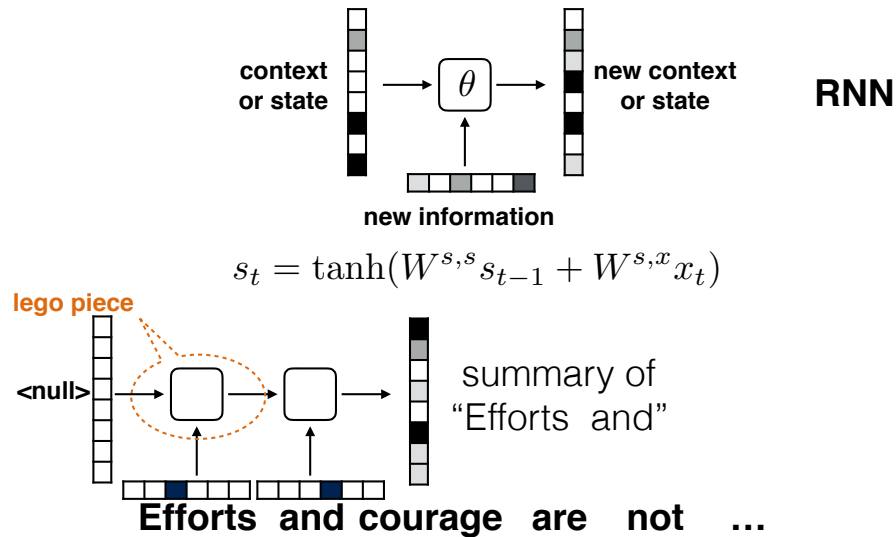
summary of
"Efforts"

Efforts and courage are not ...



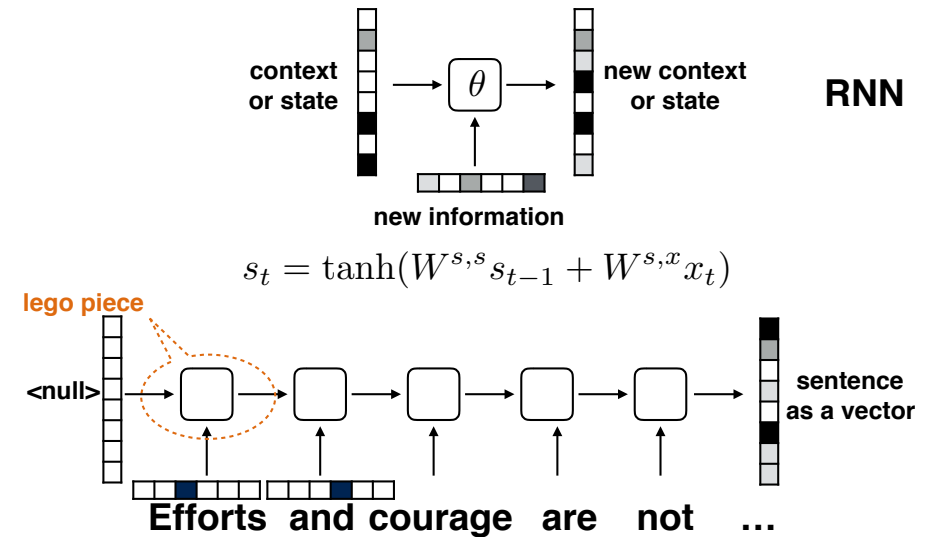
Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance



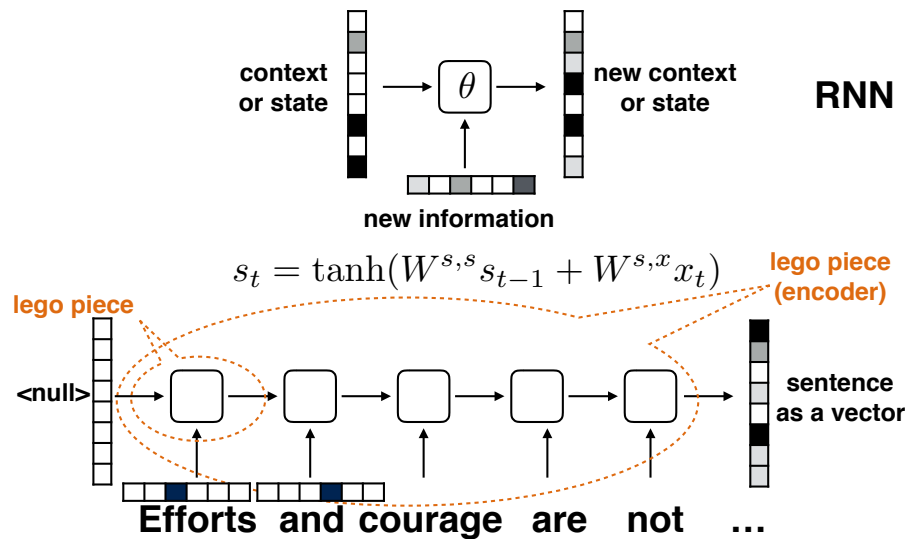
Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance




Example: encoding sentences

- Easy to introduce adjustable "lego pieces" and optimize them for end-to-end performance



Example: encoding sentences


- There are three differences between the encoder (unfolded RNN) and a standard feed-forward architecture
 - input is received at each layer (per word), not just at the beginning as in a typical feed-forward network
 - the number of layers varies, and depends on the length of the sentence
 - parameters of each layer (representing an application of an RNN) are shared (same RNN at each step)





What's in the box?

- › We can make the RNN more sophisticated...




$$s_t = \tanh(W^{s,s}s_{t-1} + W^{s,x}x_t)$$



What's in the box?

- › We can make the RNN more sophisticated...




$$g_t = \text{sigmoid}(W^{g,s}s_{t-1} + W^{g,x}x_t)$$

$$s_t = (1 - g_t) \odot s_{t-1} + g_t \odot \tanh(W^{s,s}s_{t-1} + W^{s,x}x_t)$$



What's in the box?

- › We can make the RNN more sophisticated...



$$f_t = \text{sigmoid}(W^{f,h}h_{t-1} + W^{f,x}x_t) \text{ forget gate}$$

$$i_t = \text{sigmoid}(W^{i,h}h_{t-1} + W^{i,x}x_t) \text{ input gate}$$

$$o_t = \text{sigmoid}(W^{o,h}h_{t-1} + W^{o,x}x_t) \text{ output gate}$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tanh(W^{c,h}h_{t-1} + W^{c,x}x_t) \text{ memory cell}$$

$$h_t = o_t \odot \tanh(c_t) \text{ visible state}$$

Key things

- › Neural networks for sequences: encoding
- › RNNs, unfolded
 - state evolution, gates
 - relation to feed-forward neural networks
 - back-propagation (conceptually)
- › Issues: vanishing/exploding gradient
- › LSTM (operationally)

1.

Unit 1 Lecture 8: Introduction to Machine Learning

Photo portrait of John F. Kennedy

Slides: #11

Object Source / URL: https://commons.wikimedia.org/wiki/File:John_F._Kennedy_White_House_photo_portrait_looking_up.jpg

Citation/Attribution: This file is a work of an employee of the Executive Office of the President of the United States, taken or made as part of that person's official duties.

As a work of the U.S. federal government, it is in the public domain.

2.

Unit 1 Lecture 8: Introduction to Machine Learning

John F Kennedy speech on May 25th, 1961

Slides: #11

Object Source / URL: <https://www.defense.gov/Explore/Spotlight/Apollo-11/>

Citation/Attribution: Image from defense.gov. This work is in the public domain